


Discussing

Instant Neural Graphics Primitives with a Multiresolution Hash Encoding

By Müller et.al, 2022




Oct 2022
Moritz Hambach and Rishabh Raj

Outline


- 1) understand idea
- 2) understand details and implementation

WHAT'S A NEURAL GRAPHICS PRIMITIVE?

An object represented by queries to a neural network!




EXAMPLE: IMAGE




“What’s the pixel color at X ?”

EXAMPLE: IMAGE




“What’s the pixel color at X?”

EXAMPLE: SIGNED DISTANCE FUNCTION




“How far away from surface is X ?”

EXAMPLE: SIGNED DISTANCE FUNCTION




“How far away from surface is \mathbf{X} ?”

EXAMPLE: RADIANCE & DENSITY FIELD (NERF)




“How much stuff is at X and what color does it have when viewed from d ?”

WHY NEURAL NETWORKS?



versus




- Impractical memory usage
- Rigid; limited to 3D
- But really fast

- Trades memory for compute
- Mutable; n-dim. inputs
- ~~- Reputation for being slow~~
- Competitively fast!

3 PILLARS OF INSTANT NEURAL GRAPHICS


Multiplicative speed-ups >1000x!

Rendering/training algorithm




- Task-specific GPU implementation
- **10-100x fewer steps than naïve tensor-based approach**

Small neural network



- Fully fused implementation
- **5-10x faster than TensorFlow**


“Good” input encoding



- Multiresolution hash encoding
- **Better speed-vs-quality tradeoff than prior work**
- **Task agnostic**



Rendering / training algorithm

- Depends on task
- NeRF:
 - query rays mostly near surfaces
 - step size grows linear with distance
 - stop marching once almost all light is absorbed
 - keep coarse occupancy grid to allow skipping empty space

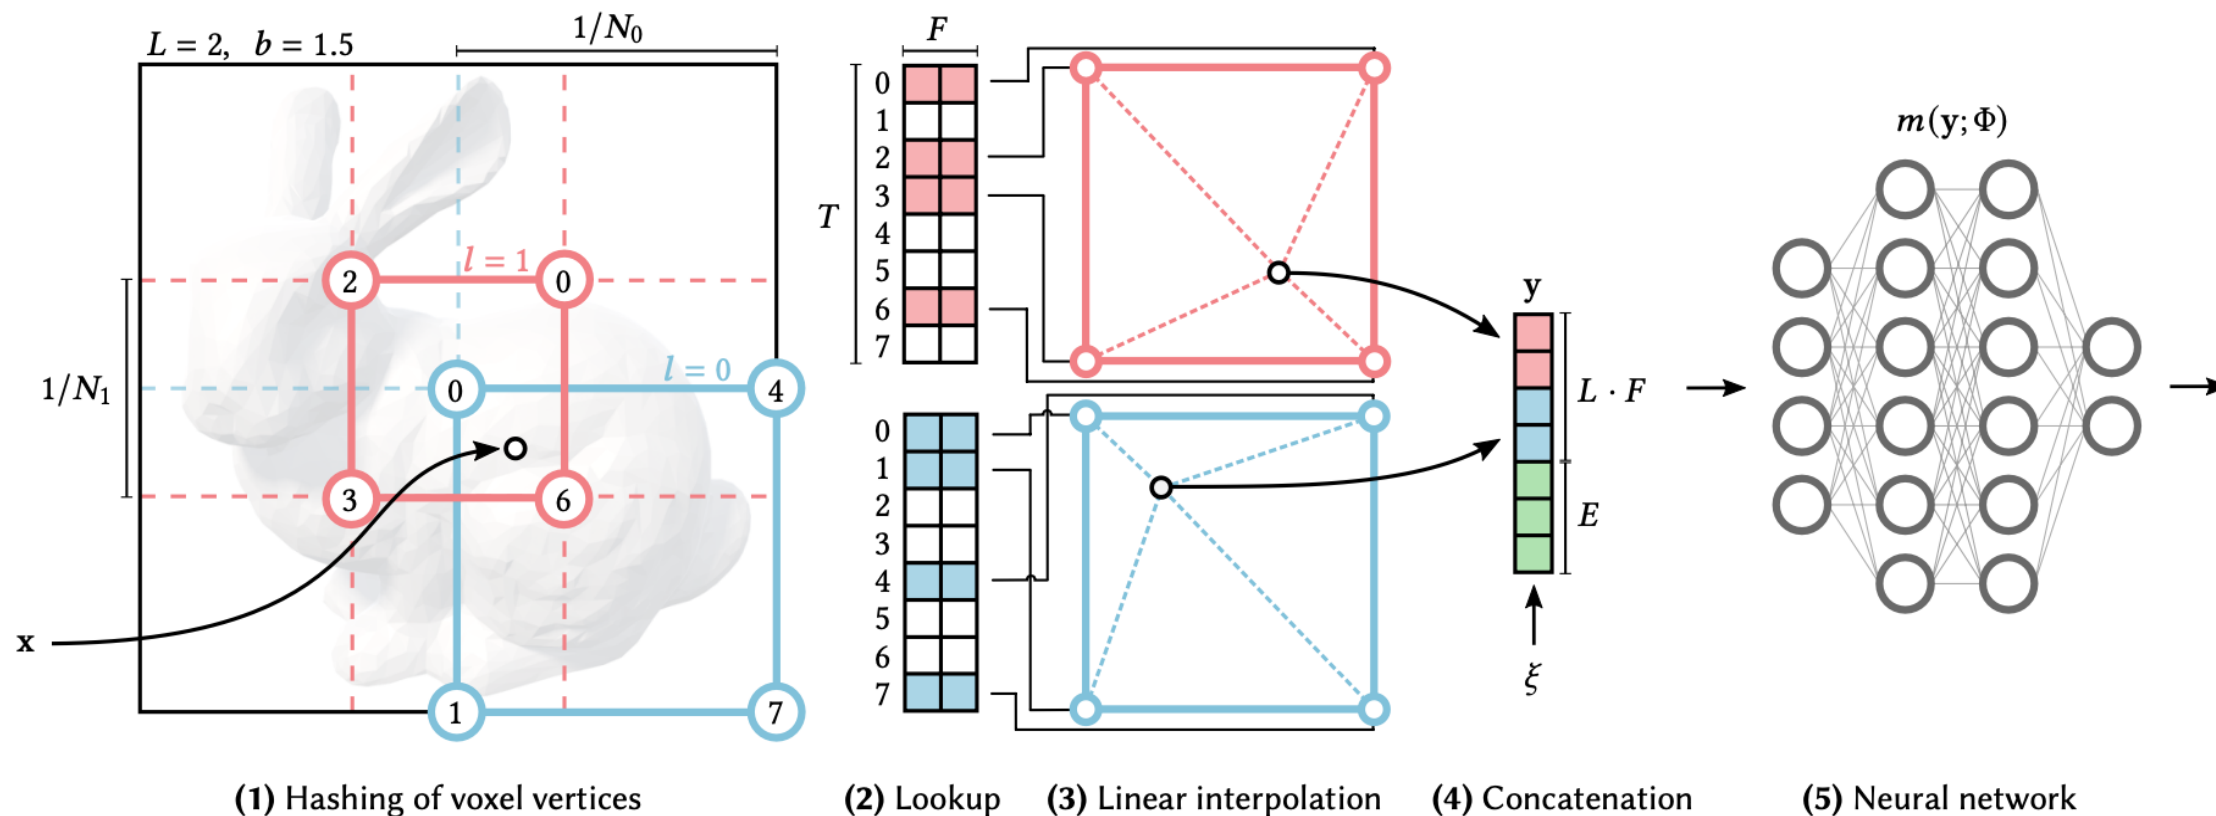


Small + fully fused Neural Networks


- MLP with only $M=64$ Neurons, $L=2..4$ hidden layers, no biases, ReLU activation
- Run on a single CUDA kernel (avoid memory traffic)
- compute is $O(M^2)$, traffic is $O(M)$, but can dominate for small M



Multi-resolution hash encoding




Multi-resolution hash encoding



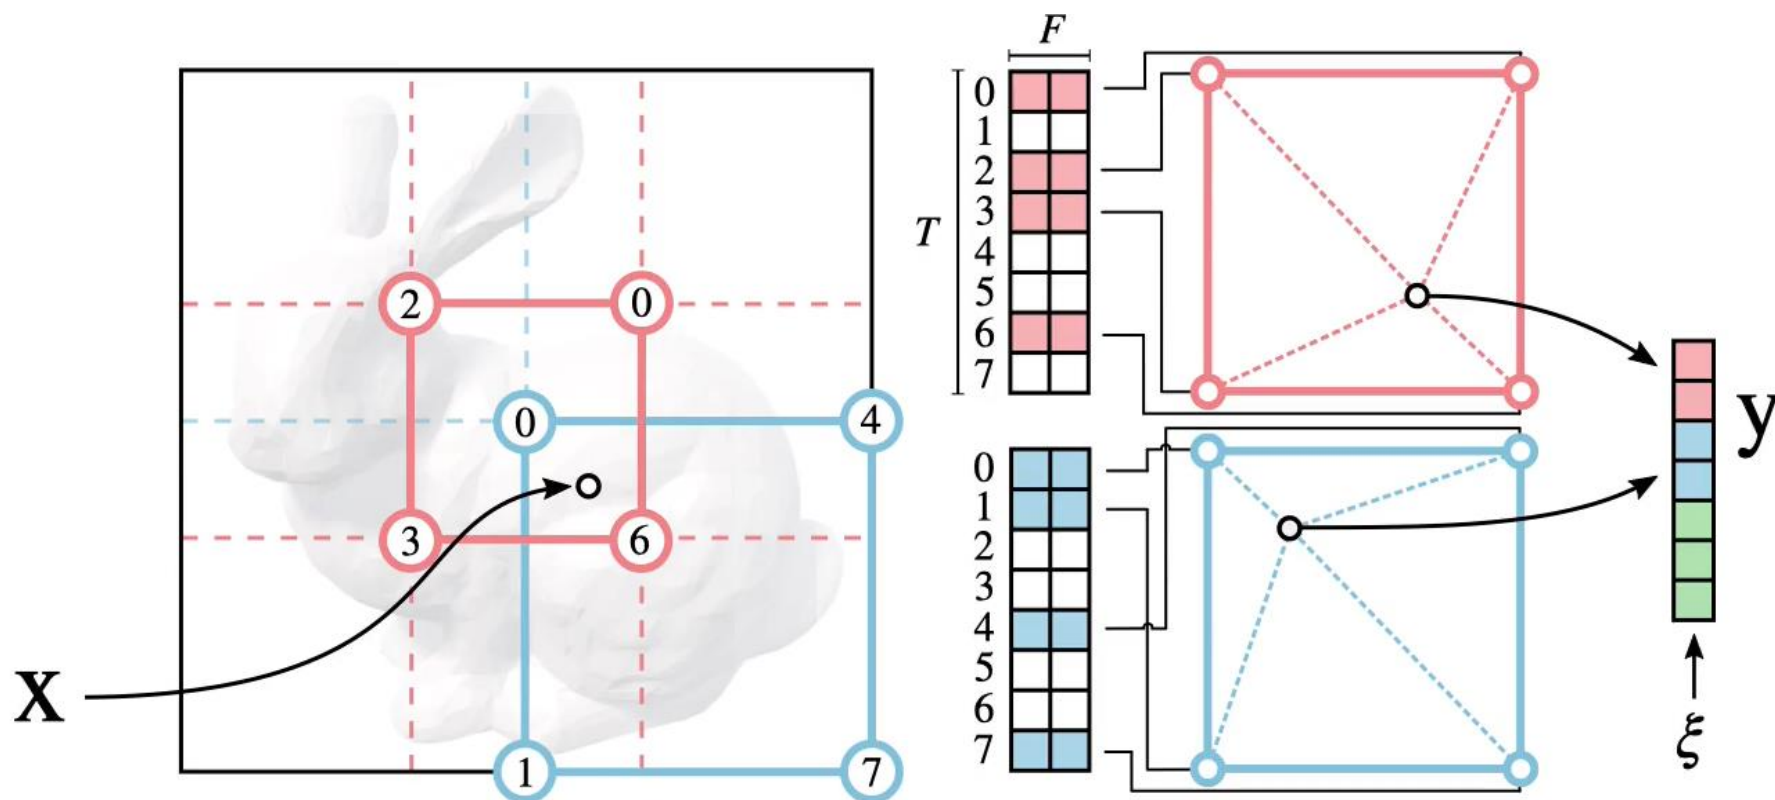
- Multi-resolution
 - Automatic level of detail

Multi-resolution hash encoding




- Multi-resolution
 - Automatic level of detail
- Hashing
 - Task agnostic
 - Table size T controls quality vs memory

Multi-resolution hash encoding



- Multi-resolution
 - Automatic level of detail
- Hashing
 - Task agnostic
 - Table size T controls quality vs memory
- Linear interpolation to enable continuity and differentiability

Multi-resolution hash encoding



Demo!

Questions?